

# MERCURY DEVELOPER'S MEETING: ANL ACTIVITIES

**PHIL CARNS** [carns@mcs.anl.gov](mailto:carns@mcs.anl.gov)

**JOHN JENKINS** [jenkins@mcs.anl.gov](mailto:jenkins@mcs.anl.gov)

Argonne National Laboratory

July 14, 2016

# ENABLING DATA SERVICES (MOCHI)

ROB ROSS, PHILIP CARNS, KEVIN HARMS,  
JOHN JENKINS, AND SHANE SNYDER

Argonne National Laboratory

GARTH GIBSON, CHUCK CRANOR, AND  
QING ZHENG

Carnegie Mellon University

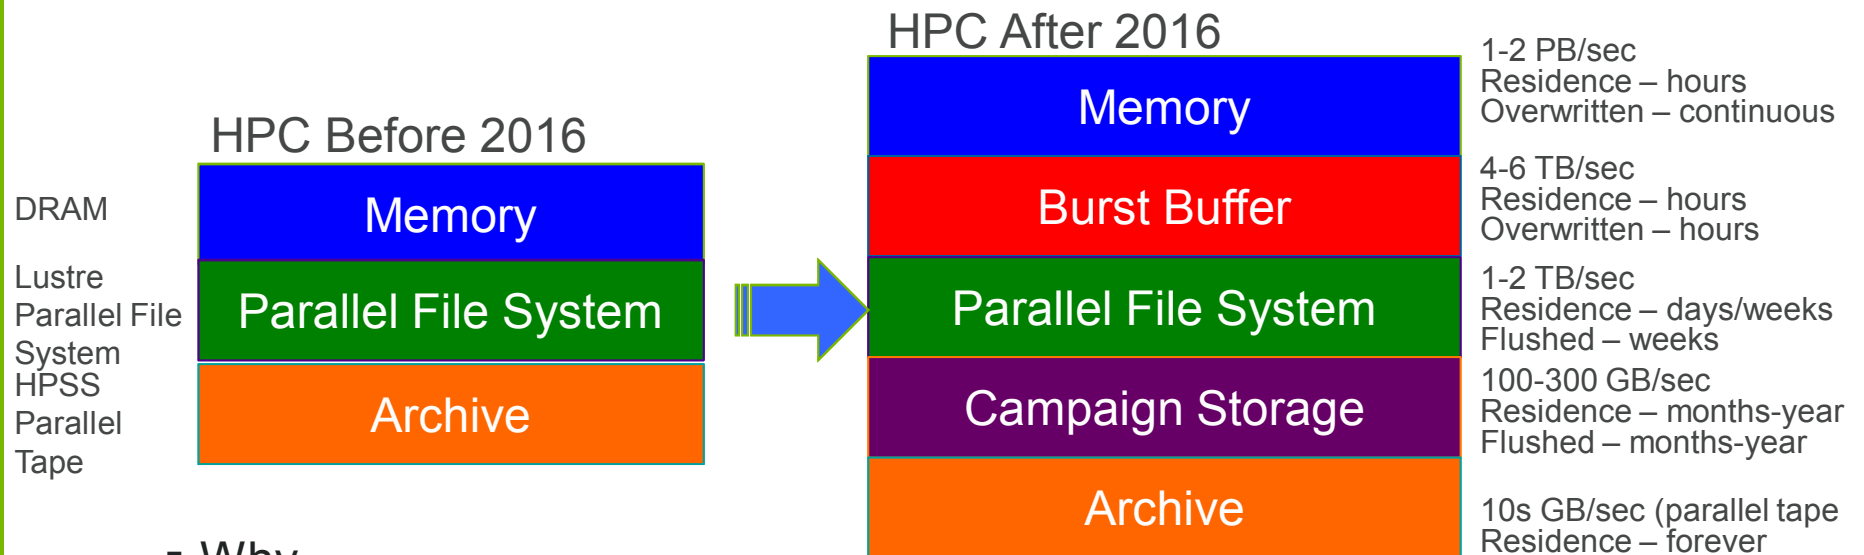
JEROME SOUMAGNE AND JOE LEE

The HDF Group

GALEN SHIPMAN AND BRAD SETTLEMYER

Los Alamos National Laboratory

# ARCHITECTURAL TRENDS: MORE STORAGE/MEMORY LAYERS...



## ■ Why

- BB: Economics (disk bw/iops too expensive)
- PFS: Maturity and BB capacity too small
- Campaign: Economics (tape bw too expensive)
- Archive: Maturity and we really do need a “forever”

Slide from Gary Grider (LANL).

# WHAT COMES NEXT?

- Assumptions
  - New layers in storage hierarchy, lower latencies
  - Storage resources will be highly contended for
  - No “holy grail” emerges that solves everyone’s problems
- Alternative model to the “PFS for data management”
  - Multiple services employed for different classes of data
  - Specialization for scalability/efficiency/productivity
  - In some cases, co-design with applications

Specialized data services  
are already here!

	Provisioning	Comm.	Local Storage	Fault Mgmt. and Group Membership	Security
<b>ADLB</b> <i>Data store and pub/sub.</i>	MPI ranks	MPI	RAM	N/A	N/A
<b>DataSpaces</b> <i>Data store and pub/sub.</i>	Indep. job	Dart	RAM (SSD)	Under devel.	N/A
<b>DataWarp</b> <i>Burst Buffer mgmt.</i>	Admin./ sched.	DVS/ Inet	XFS, SSD	Ext. monitor	Kernel, Inet
<b>FTI</b> <i>Checkpoint/restart mgmt.</i>	MPI ranks	MPI	RAM, SSD	N/A	N/A
<b>Kelpie</b> <i>Dist. in-mem. key/val store</i>	MPI ranks	Nessie	RAM (Object)	N/A	Obfusc. IDs
<b>SPINDLE</b> <i>Exec. and library mgmt.</i>	Launch MON	TCP	RAMdisk	N/A	Shared secret

# OUR GOAL

## Enable composition of data services for DOE science and systems

- Application-driven
  - Identify and match to science needs
  - Traditional data roles (e.g., checkpoint, data migration)
  - New roles (e.g., equation of state/opacity databases)
- Composition
  - Develop/adapt building blocks
    - Communication
    - Concurrency
    - Local Storage
    - Resilience
    - Authentication/Authorization
- Enable rapid development of specialized services
- ***Don't built new services from scratch every time***



# SERVICE COMPONENTS AND MERCURY INTEGRATION

ROBERT DASS, PHILIP CARNIS, KEVIN THOMAS, AND JOHN JENKINS, AND SHANE SNYDER Argonne National Laboratory

CAROL GIBSON, CHUCK CRANCE, AND JING ZHENG Carnegie Mellon University

JEROME SOUMAGNE AND JOE LEE The HDF Group

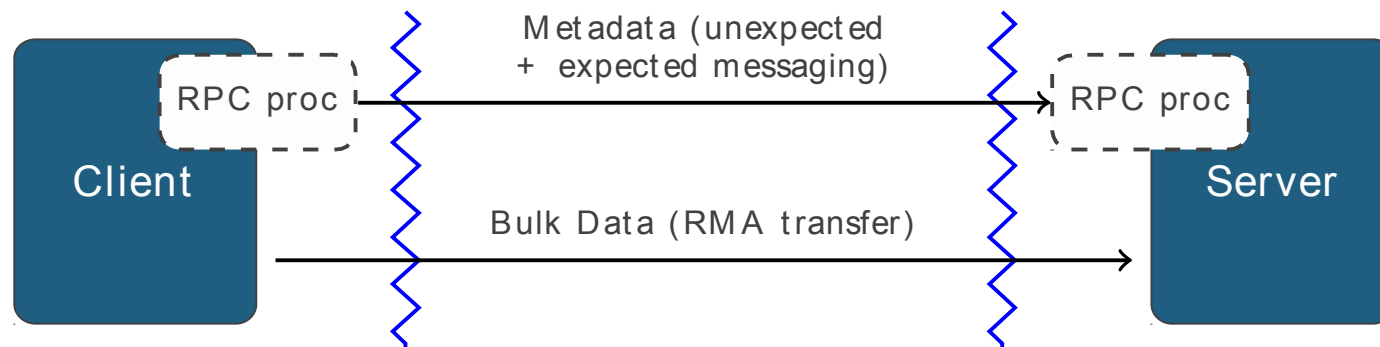
GALEN SHIPMAN AND BRAD SETTLEMYER Los Alamos National Laboratory

# COMMUNICATION: MERCURY

<https://mercury-hpc.github.io/>

The Mercury RPC system is the core building block of all of the services we are building in the Mochi project.

- Portable across systems and network technologies
- Builds on lessons learned from IOFSL, Nessie, Inet, and others
- Efficient bulk data movement to complement control messages





# PROGRAMMING MODELS FOR MERCURY SERVICES

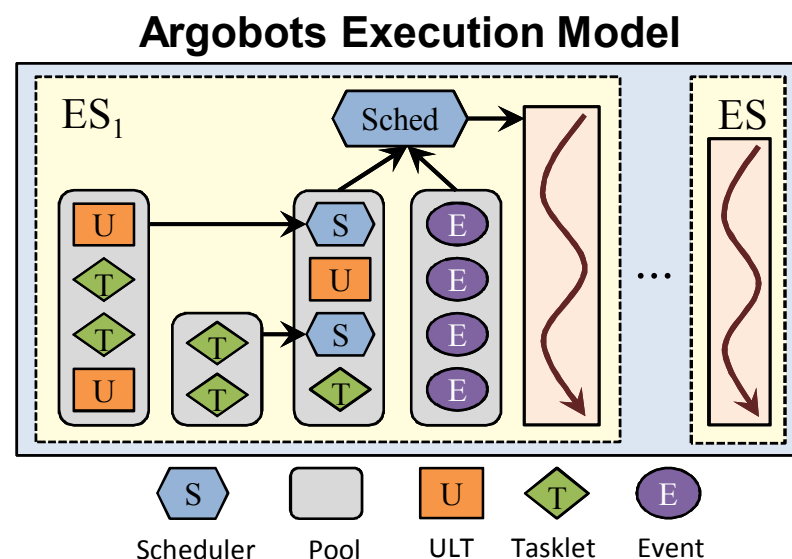
## Concurrency, low latency, and ease of use

- The Mercury API uses an event-driven callback model that performs well, but we need more than that to build flexible services rapidly
  - Programmability of callback continuation (i.e. avoid stack-ripping)
  - Assignment of work to CPU cores
  - Scheduling for extremely high concurrency
- We've elected to combine Mercury with ***user-level threading***
  - Similar to co-routines, green threads, fibres, etc.
  - Hide callbacks from service developer
  - Combine latency benefits of Mercury's event model with programmability of conventional thread model

# CONCURRENCY: ARGOBOTS

## Argobots is a lightweight threading/tasking framework

- Features relevant to I/O services:
  - Flexible mapping of work to hardware resources  
(*which ULTs can run on which cores*)
  - Ability to delegate service work with fine granularity across those resources  
(*lightweight tasklets and ULTs*)
  - Modular, user-definable scheduling
- We developed asynchronous bindings to:
  - Mercury
  - LevelDB
  - POSIX I/O
- Working with Argobots team to identify needed functionality (e.g., idling) for our use cases



# MARGO

## Binding Mercury to Argobots

- Margo is a support library that provides Argobots-aware bindings to Mercury
  - Mercury operations are normally expressed with a post/callback model
  - Margo instead provides blocking functions that will suspend and resume the calling ULT automatically
- Internal progress loop drives Mercury progress and callbacks
  - Can run on dedicated core/thread or share resources of caller
- Incoming RPC handlers are executed as new user level threads
- “\_timed()” versions of each function can be used to specify timeouts for cancellation/completion

# WHAT DOES THE CODE LOOK LIKE?

## A few Margo examples

- Macros are used to define wrappers for RPC handlers that will be launched as user-level-threads  
<https://xgitlab.cels.anl.gov/sds/margo/blob/master/examples/my-rpc.c#L111>
- Services code (both clients and servers) issue Mercury calls as blocking operations  
<https://xgitlab.cels.anl.gov/sds/margo/blob/master/examples/my-rpc.c#L58>
- Explicit concurrency is achieved by creating and joining user-level threads  
<https://xgitlab.cels.anl.gov/sds/margo/blob/master/examples/client.c#L137>
- Boilerplate: after starting Mercury library, set up Argobots and tell Margo what resources to use for driving progress and running RPC handlers  
<https://xgitlab.cels.anl.gov/sds/margo/blob/master/examples/server.c#L53>



# HIGH-LEVEL TUNING OBSERVATIONS

## Observations from Argobots/Margo/Mercury/CCI stack

- CPU usage
  - Trade-off between latency and how frequently you are willing to busy-spin on a CPU core
  - CCI level
    - Reduce CPU usage by using CCI poll feature  
-DNA\_CCI\_USE\_POLL:BOOL=ON), recently upstreamed in Mercury
  - Argobots level
    - Reduce CPU usage with abt-snoozer scheduler (optional add-on for use with Margo and other support libraries that block on I/O)
- Bandwidth
  - So far so good for large transfers

# HIGH-LEVEL TUNING OBSERVATIONS (CONT)

## Observations from Argobots/Margo/Mercury/CCI stack

- Latency
  - Pay close attention to callback scheduling and context switching points in progress loop (Margo in our case)
  - Disable checksumming in Mercury  
(-DMERCURY\_USE\_CHECKSUMS:BOOL=OFF)
  - Tcmalloc reduces memory management latency
- Ongoing
  - Isolating sources of variability
  - Investigating memory allocation and reuse
  - Testing more scenarios

*We need to standardize on how to measure and tune more of these things...*

# MERCURY DEVELOPER'S MEETING: ANL EXAMPLES AND BENCHMARKING

JOHN JENKINS [jenkins@mcs.anl.gov](mailto:jenkins@mcs.anl.gov)

PHIL CARNS [carns@mcs.anl.gov](mailto:carns@mcs.anl.gov)

Argonne National Laboratory

July 14, 2016

# DEMONSTRATING OUR APPROACH



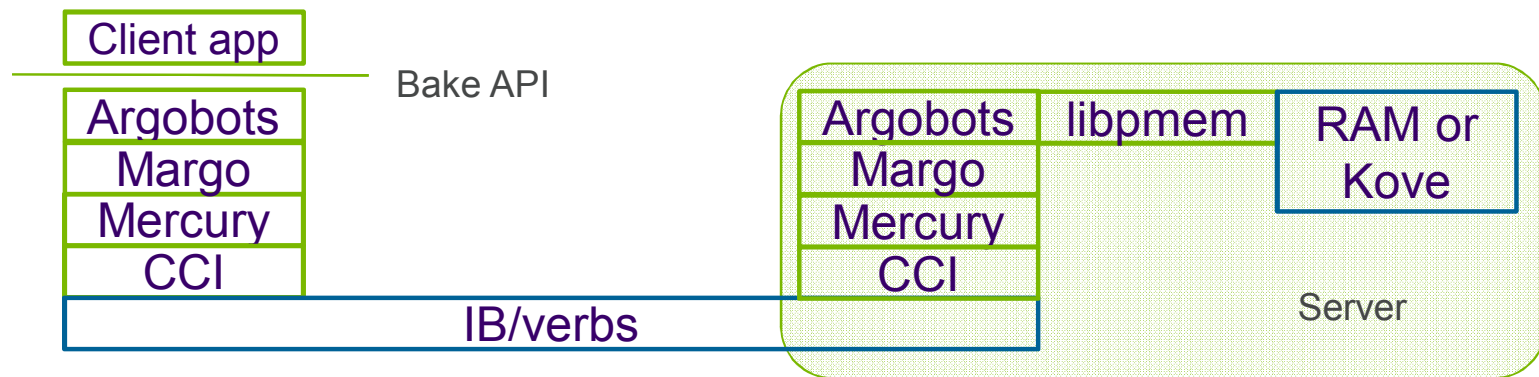
# DEPLOYMENT SCENARIOS

## When, where, and how to run composable data services

- Use cases will depend on both the system and the application, so we need to be flexible
- Examples:
  - Co-located alongside (or even within) application processes
  - Set-aside nodes in a single job
  - Co-scheduled across jobs
  - Persistent (or semi-persistent) services
- Storage may reside on compute nodes, burst buffer nodes, or off-system
- Duration may be transient, on-demand, or long-running
- Some of these scenarios will rely on resource management integration

# BUILDING BLOCKS FOR STORAGE SERVICES

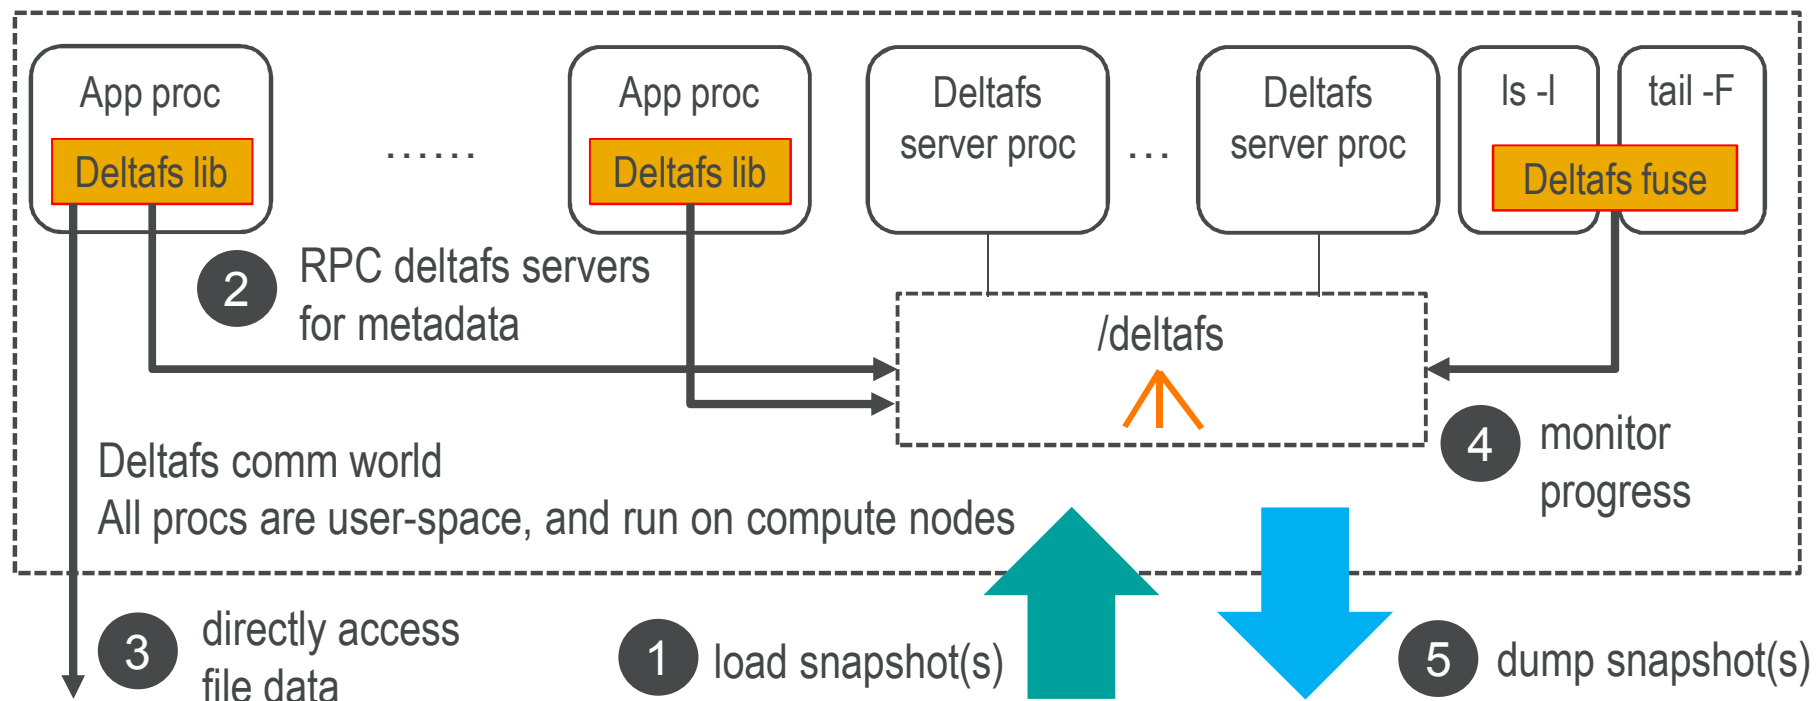
- In-progress low-level storage service, “BAKE”, to serve as a basis for future data services



- Looks like a lot of components
- ... but data copies and context switches are rare
- RDMA transfer directly from client app address space to storage device that has been mapped into server address space
- Argobots/Margo/Mercury handle concurrency without any OS threads

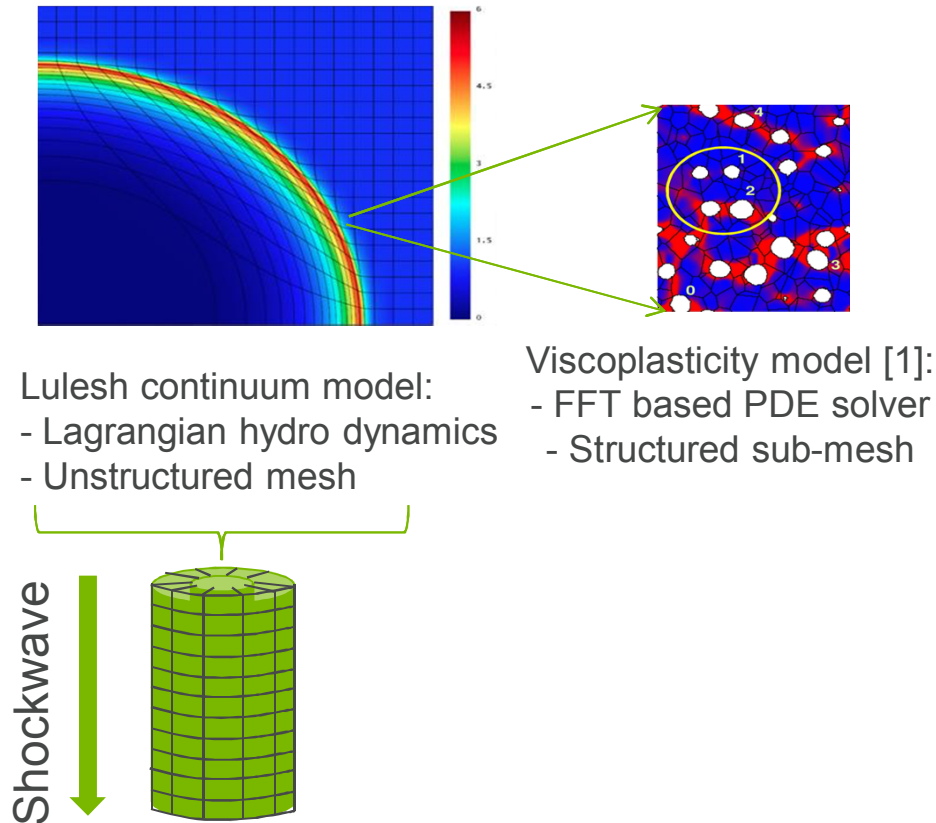
# TRANSIENT FILE SYSTEM VIEWS: DELTAFS

Supporting legacy POSIX I/O in a scalable way.



Credit: Qing Liu, CMU

# CO-DESIGNING COUPLED MODELS



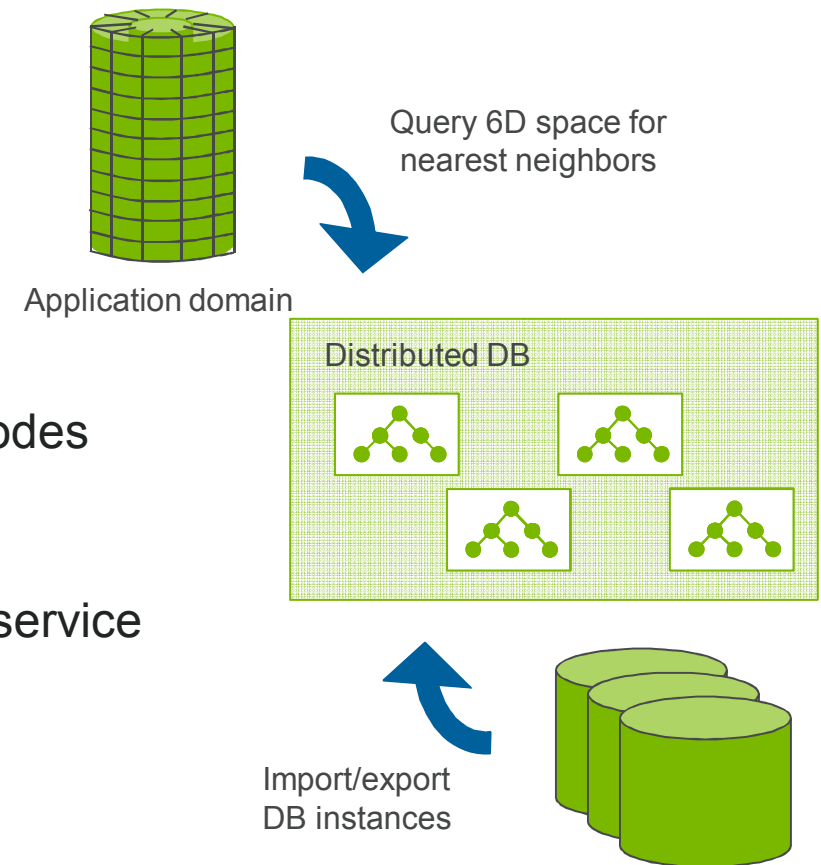
- Future applications are exploring the use of multi-scale modeling
- As an example: Loosely coupling continuum scale models with more realistic constitutive/response properties
  - e.g., Lulesh from ExMatEx
- Fine scale model results can be cached and new values interpolated from similar prior model calculations

R. Lebensohn et al, Modeling void growth in polycrystalline materials, Acta Materialia, <http://dx.doi.org/10.1016/j.actamat.2013.08.004>.



# CO-DESIGNING A FINE SCALE MODEL DATABASE

- Goals
  - Minimize fine scale model executions
  - Minimize query/response time
  - Load balance DB distribution
- Approach
  - Distributed approx. nearest-neighbor query
  - Partitioned spatial data structures across nodes
  - Import/export to persistent store
- Status
  - Mercury-based, centralized in-memory DB service
  - Investigating distributed, incremental nearest-neighbor indexing

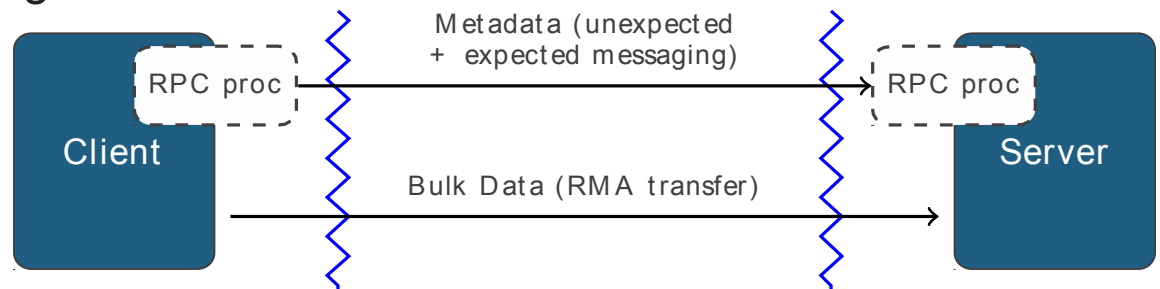


# BENCHMARKING

# MICROBENCHMARKS

Initial code at <https://github.com/mercury-hpc/mercury-benchmarks>

- RPC round-trip latency
- Bulk transfer rate
- Concurrent RPCs/bulks
- Overheads (not yet directly measured)
  - RPC data (un)marshalling
  - HG/NA work queue management
  - NA → plugin translation



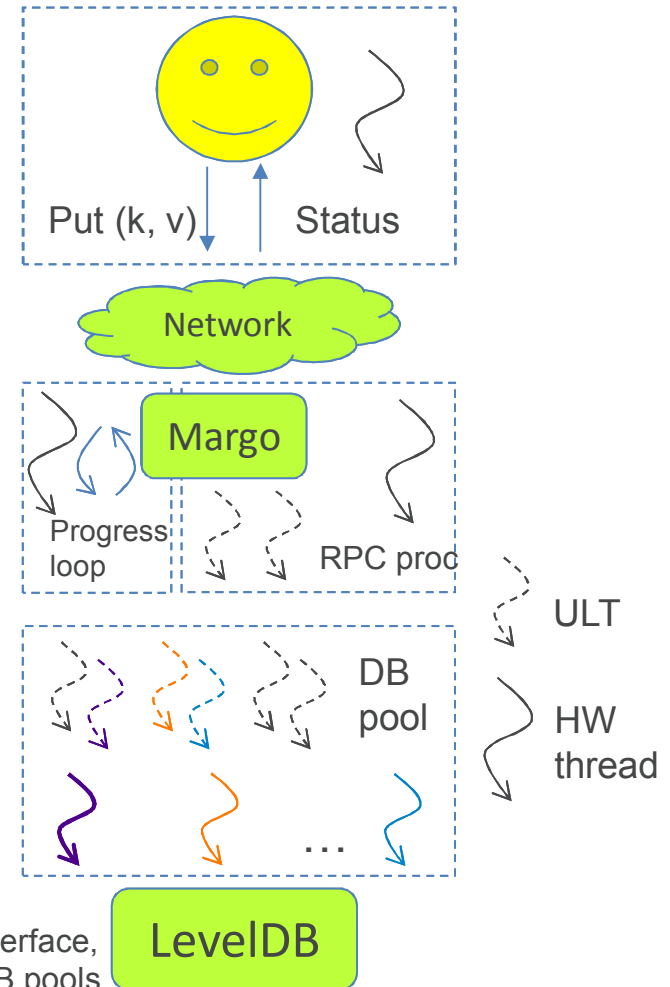
*Network Abstraction Layer*

# BENCHMARKING COMMON DESIGN PATTERNS

## Capturing Mercury's surrounding ecosystem

- Example: LevelDB wrapper
- Mixed RPC+bulk workflow
  - k/v pair represented as a bulk handle
  - Server receives “put” RPC, RDMA reads using k/v handle
- Multithreading scenarios
  - HG progress, DB op setup / dispatch
- Interaction with other layers
  - Argobots ULT creation, context switching, thread handoff

**NOTE:** LevelDB doesn't have an async interface, otherwise could unify the RPC proc and DB pools





# CONCLUDING REMARKS

# TOPICS FOR AFTERNOON DISCUSSION

- Multi-user environments (secure messaging)
- Group membership (dynamic, fault tolerant service participation)
- Transport roadmaps (libfabric)
  - currently using BMI from PVFS, CCI as prototyping vehicles
- Standardizing set of benchmarks?
- Leveraging LNET experience

# RESOURCES

- Mercury: <https://github.com/mercury-hpc/mercury>
- Mercury benchmarks: <https://github.com/mercury-hpc/mercury-benchmarks>
- Argobots: <https://collab.cels.anl.gov/display/ARGOBOTS/Argobots+Home>
- Margo: <https://xgitlab.cels.anl.gov/sds/margo>
- abt-snoozer: <https://xgitlab.cels.anl.gov/sds/abt-snoozer>



**THIS WORK IS SUPPORTED BY THE DIRECTOR, OFFICE OF  
ADVANCED SCIENTIFIC COMPUTING RESEARCH, OFFICE OF  
SCIENCE, OF THE U.S. DEPARTMENT OF ENERGY UNDER  
CONTRACT NO. DE-AC02-06CH11357.**