



UNIVERSITY OF OREGON



Scalable Observation System for Scientific Workflows

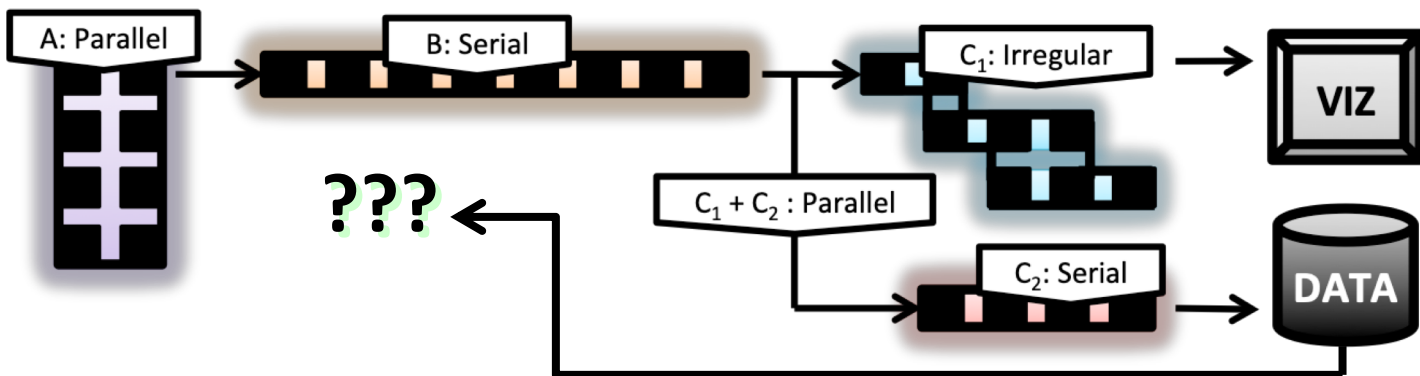
Overview, Usage, Resiliency, & Security

Chad Wood, Kevin Huck

SC18 BoF: Enabling Data Services for HPC, November 13th 2018

- ❖ It is reasonable to want to **use** “information” during application execution, and this is uniquely challenging
- ❖ **More than just performance data, an integrated perspective**
- ❖ Information comes from the application as well as from the environment in which the application is executing
- ❖ *Application*: Performance, problem-specific data and metadata, ...
- ❖ *Environment*: System state, resource usage, runtime properties, ...
- ❖ *Workflow*: Several components running together, sharing resources
- ❖ *Campaign*: Access to data within and across allocations

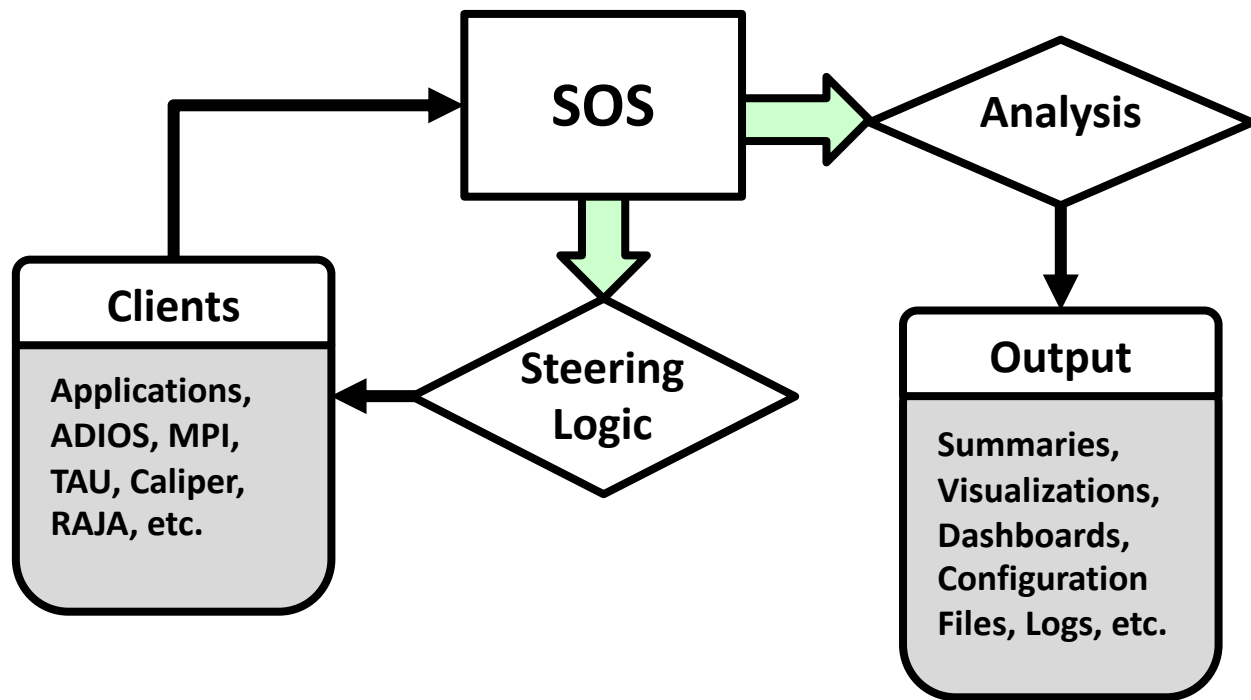
Where SOS is Needed



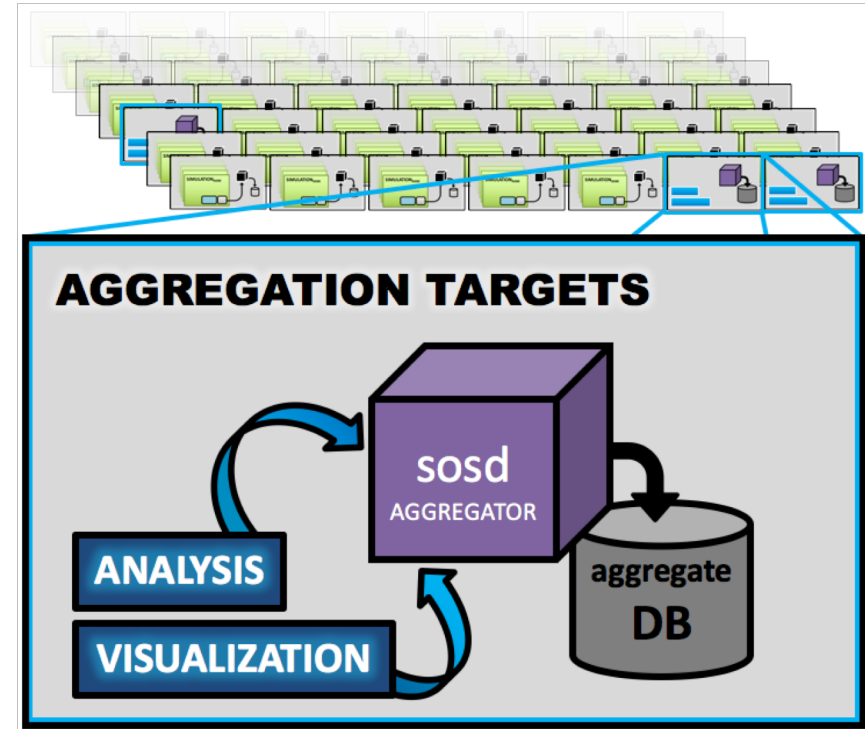
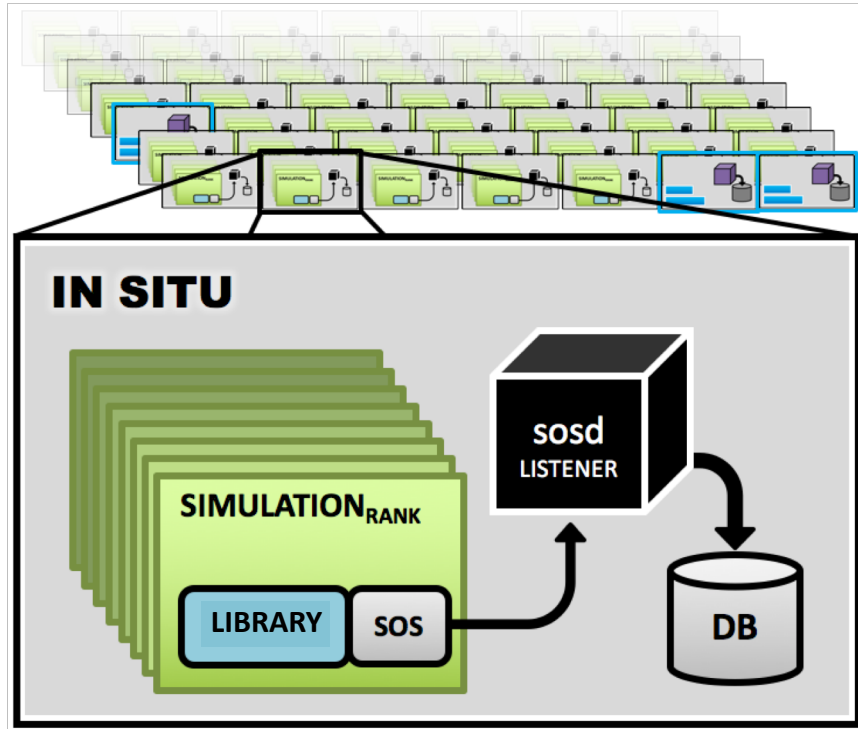
- ❖ Multiple components with *data flow*
- ❖ Applications using *MPI* and *filesystem* for coordination are sensitive to interference
- ❖ Complex interactions with *dynamic behavior*
- ❖ Wherever *variability is evident* in application behavior, hardware, or shared resource utilization during a run
- ❖ Situations where *acting on information is possible*

SOSflow functions as a hub for collecting, aggregating, and acting on a variety of information at runtime.

SOSflow's in situ (online) services work together to provide global views and online data analytics in an HPC environment.



SOS Topology : Deployment



Integration

- ❖ Runs entirely in user space
- ❖ Direct use or via tools:
 - ❑ **TAU**: Asynchronous periodic output to SOS
 - ❑ **Caliper**: Enable SOS service to receive event-driven output
 - ❑ **Alpine**: Geometry extraction, projection
 - ❑ **Apollo & RAJA**: ML for portable performance
 - ❑ **Python, C++, C** APIs
- ❖ **Challenges**:
 - ❑ Coordinating start-up and shutdown w/apps

Resiliency

- ❖ SOS is highly configurable to many scenarios, with trade/offs in data persistence, analysis support, and shared resource use facilitated by:
 - ❑ File-based SQL DB
 - ❑ Memory-only SQL DB
 - ❑ Circular Cache
 - ❑ Flexible aggregation and analysis topology
- ❖ **Challenges**:
 - ❖ Drowning by firehose
 - ❖ Client-library response to service failure

Security

- ❖ How much is needed?
 - ❑ Edge security covers user space (for now)
 - ❑ Dashboards, Jupyter
 - ❑ Inter-machine coordination
 - ❑ Long-term storage of collected data
 - ❑ What resources are available
- ❖ **Current solution: Munge**
 - ❑ Used by DoE for MPI
 - ❑ Practical isolation
 - ❑ Offers many more features than we use